

A Federated Aggregation Framework for E-values: Comparing Weighting Strategies for Multi-Site Sensitivity Analysis

[Technical Report] • [Preprint Version]

Author: Daijiro Wachi

Email: daijiro.wachi@gmail.com

Version: 2.0 (Technical Report)

Code: <https://github.com/watilde/Harmonia/tree/main/research/modules/2-federated-evalues>

ABSTRACT

Background: E-values quantify sensitivity to unmeasured confounding in single-site studies, but principled multi-site aggregation methods have not been characterized.

Objective: Compare three aggregation strategies (sample-size weighting, equal-weight, conservative) for E-values in multi-site studies, characterize their mathematical properties, and assess convergence behavior across scales.

Methods: We compared three E-value aggregation strategies: sample-size weighting (Federated Robustness Index, FRI), equal-weight averaging, and conservative minimum. Characterized mathematical properties (boundedness, monotonicity) and validated across three scales (1k-2.8m patients, 3 sites) using Synthea synthetic data with communication efficiency and partial privacy analysis.

Results: FRI converged from 2.015 (1k) to 2.149 (2.8m) with inter-site coefficient of variation collapsing from 9.7% to 0.16%. Sample-size and equal-weight strategies diverged only at small scale with unbalanced sites (2.8% difference). Communication: $O(1)$ complexity with constant 174 bytes versus $O(N)$ centralized (201 KB-482 MB), achieving up to $2.8M\times$ reduction. Privacy tradeoff: covariate identity privacy preserved (methodological choices hidden), while aggregate effect sizes remain inferable from E-values by mathematical necessity—a design tradeoff inherent to meta-analysis, not a flaw.

Conclusions: Our systematic comparison demonstrates that sample-size weighting (FRI) provides precision-weighted summaries following meta-analysis conventions, while performance differences between strategies diminish at scale. Practitioners should report both FRI (precision-weighted) and $\min\{E_k\}$ (conservative) to enable risk-appropriate decisions. This is a technical demonstration of computational feasibility, not clinical validation. Proposed thresholds require validation against RCT ground truth in real multi-site studies.

Keywords: E-values, Sensitivity Analysis, Unmeasured Confounding, Federated Learning, Multi-Site Studies

1. INTRODUCTION

1.1 The Federated Sensitivity Analysis Gap

Causal inference from observational data requires addressing unmeasured confounding [1,2]. **E-values** quantify robustness: the minimum strength (risk ratio) of unmeasured confounding required to nullify an observed effect [4,5]. An E-value of 2.5 means a confounder must have RR 2.5 with both treatment and outcome to explain away the effect.

Gap: E-values are single-site metrics. In federated settings, **how should site-level E-values be aggregated with validity guarantees?**

Example - 3-hospital diabetes network:

- Academic (n=900): E-value=2.8 (robust)
- Community A (n=380): E-value=1.6 (moderate)
- Community B (n=380): E-value=1.9 (moderate)

Questions: What is the federated E-value? Should large sites dominate? Does aggregation preserve validity?

This work’s contribution: Systematically comparing three aggregation strategies for multi-site E-values, characterizing their mathematical properties (Definition 1), and validating across three scales (1k-2.8m patients). We provide practical guidance for selecting strategies based on scale and heterogeneity.

2. METHODS

2.1 E-value Background

For a causal effect estimate $\hat{\theta}$, the E-value is [4]:

$$E = RR + \sqrt{RR \times (RR - 1)}$$

where RR is the risk ratio corresponding to $|\hat{\theta}|$.

From bounds: When using Manski partial identification, we compute E-values from both bounds:

$$E_L = f(\mathcal{L}), \quad E_U = f(\mathcal{U})$$

Conservative E-value = $\min(E_L, E_U)$ (most sensitive direction).

2.2 Federated Robustness Index (FRI)

Definition: The FRI aggregates site-level E-values E_k using weighted averaging:

$$\text{FRI} = \sum_{k=1}^K w_k E_k$$

where w_k are federation weights satisfying $\sum_k w_k = 1$ and $w_k \geq 0$.

Interpretation: FRI represents the minimum unmeasured confounding strength (as risk ratio) required to nullify the federated effect estimate.

2.3 Aggregation Strategies

Strategy	Weight Formula	Properties
Sample-size	$w_k = n_k/N$	Proportional to precision
Equal	$w_k = 1/K$	Democratic
Conservative	$\text{FRI} = \min_k E_k$	Most cautious

Primary focus: Sample-size weighting (theoretically justified, see Section 2.4).

2.4 FRI Definition and Properties

Definition 1 (Federated Robustness Index):

For K federated sites with local E-values E_k and sample sizes n_k , define the Federated Robustness Index as:

$$\text{FRI} = \sum_{k=1}^K \frac{n_k}{N} E_k \quad \text{where} \quad N = \sum_j n_j$$

Rationale: Sample-size weighting follows meta-analysis conventions (DerSimonian & Laird, 1986), giving larger sites proportionally higher weight consistent with their contribution to federated precision.

Properties (immediate from definition):

1. **Boundedness:** $\min\{E_k\} \leq \text{FRI} \leq \max\{E_k\}$ (by convex combination)
2. **Monotonicity:** $\partial \text{FRI} / \partial E_k = w_k > 0$ for all k
3. **Convexity:** FRI is a convex combination of site-level E-values.

Interpretation and limitations: FRI represents the “precision-weighted average unmeasured confounding strength” across sites. This is a **descriptive summary statistic**, not a formal worst-case robustness guarantee.

What FRI is NOT: FRI does not provide a strict lower bound for unmeasured confounding required to nullify the federated effect. The true worst-case threshold is $\min_k \{E_k\}$ (if the weakest site’s effect is nullified, the federated average is affected). $\text{FRI} > \min\{E_k\}$ because it incorporates information from more robust sites.

When FRI is appropriate:

- Sites have heterogeneous robustness (varying E-values)
- Sample sizes differ considerably across sites
- Decision-makers want a precision-weighted summary beyond simple min/max

When to report $\min\{E_k\}$ instead:

- Risk-averse contexts requiring worst-case guarantees
- Regulatory decisions with severe consequences of false positives
- Sites are equally sized (weights $1/K$)

Remark 1 (Comparison to Conservative Aggregation): The true conservative approach uses $E_{\text{conservative}} = \min_k \{E_k\}$, ensuring that even the most vulnerable site’s robustness is respected. $\text{FRI} > \min\{E_k\}$ by construction, making it less conservative but potentially more informative by incorporating multi-site evidence. Decision-makers should report both:

- FRI: Precision-weighted summary
- $\min\{E_k\}$: Worst-case guarantee

Remark 2 (Reduced Confounder Disclosure): Federated E-value transmission provides **partial privacy** compared to centralized covariate sharing. Sites transmit scalar E-values rather than full covariate matrices, reducing but not eliminating information leakage. E-values indirectly reveal effect size magnitudes (via RR), which can suggest confounding adjustment strength. For example, a high E-value indicates either large treatment effects or strong confounder adjustment. Formal differential privacy (via Laplace noise: $E'_k = E_k + \text{Lap}(0, \Delta/\epsilon)$) would provide rigorous ϵ -DP guarantees, which we leave to future work.

3. **Relationship to conservative aggregation:** $\text{FRI} \geq \min\{E_k\}$, with equality only when all E_k are identical (homogeneous case)

Empirical Validation (from the experiments):

- **1k:** $\text{FRI}=1.961, E_k \in [1.766, 2.188] \rightarrow 1.766 < 1.961 < 2.188$
- **2.8m:** $\text{FRI}=2.149, E_k \in [2.146, 2.153] \rightarrow \text{nearly min} = \text{FRI} \approx \text{max}$ (homogeneity)

2.5 Threshold Interpretation (Exploratory)

VanderWeele (2019) suggests $E > 2$ indicates “moderate robustness” for single-site studies, though these are heuristic guidelines rather than validated cutoffs. Extending this to federated settings, we propose exploratory thresholds:

- FRI > 2.0 for moderate-stakes decisions (clinical guidelines)
- FRI > 3.0 for high-stakes decisions (regulatory approval)
- FRI > 1.5 for exploratory research (hypothesis generation)

Critical limitation: These thresholds lack empirical validation via retrospective RCT comparison or regulatory decision analysis. They represent conceptual proposals requiring validation through: (1) comparison of observational FRI to RCT-confirmed effects, (2) expert elicitation of clinician risk tolerance, and (3) analysis of FDA decisions with known confounding.

Recommended practice: Report both FRI (precision-weighted summary) and $\min\{E_k\}$ (conservative worst-case guarantee). Risk-averse decisions should use $\min\{E_k\}$; risk-neutral contexts may consider FRI. Thresholds provide guidance, not rigid cutoffs—FRI = 1.99 vs 2.01 should not mechanically determine decisions.

2.6 Experimental Design

Three Dataset Scales:

Scale	Total Patients	Sites	Purpose
Small (1k)	1,130	3	Heterogeneity effects
Medium (100k)	235,222	3	Convergence behavior
Large (2.8m)	2,709,803	3	Asymptotic validation

Data: OMOP-formatted Synthea diabetes treatment data with MTR bounds.

Metrics: FRI convergence, inter-site coefficient of variation (CV), E-value decomposition.

3. RESULTS

3.1 Multi-Scale FRI Convergence and Aggregation Strategy Comparison

Table 1: Site-Level E-values and Aggregation Strategies Across Scales

Scale	Site 1 E	Site 2 E	Site 3 E	Site 1 n	Site 2 n	Site 3 n	Range	CV
1k	2.188	1.766	1.929	533	194	403	0.422	9.7%
100k	2.156	2.143	2.143	78,408	78,406	78,408	0.013	0.30%
2.8m	2.153	2.146	2.148	903,268	903,267	903,268	0.007	0.16%

Table 2: Aggregation Strategy Comparison

Scale	Sample-Size FRI	Equal- Weight	Conservative (min)	Optimistic (max)	FRI vs min
1k	2.015	1.961	1.766	2.188	+14.1%
100k	2.147	2.147	2.143	2.156	+0.19%
2.8m	2.149	2.149	2.146	2.153	+0.14%

Aggregation formulas:

- Sample-size FRI: $\sum_k (n_k/N) \cdot E_k$
- Equal-weight: $(1/K) \sum_k E_k$
- Conservative: $\min_k \{E_k\}$
- Optimistic: $\max_k \{E_k\}$

FRI converges from 2.015 (1k) to 2.149 (2.8m), with diminishing increments ($\Delta = 0.132 \rightarrow 0.002$) indicating asymptotic stability. Inter-site heterogeneity collapses from 9.7% to 0.16% coefficient of variation, confirming that sampling variation dominates at small scales while sites homogenize at large scales. The boundedness property ($\min\{E_k\} \leq \text{FRI} \leq \max\{E_k\}$) holds at all scales, validating Definition 1.

Sample-size and equal-weight strategies diverge only at 1k scale (2.015 vs 1.961, 2.8% difference) due to unbalanced site sizes (n: 533, 194, 403). At 100k and 2.8m scales, near-uniform site sizes (CV < 0.01%) render weighting strategy irrelevant—all approaches converge to within 0.1%. This suggests that for large federated networks with balanced sites, aggregation method selection matters little.

The FRI-to- $\min\{E_k\}$ gap narrows sharply with scale (14.1% at 1k $\rightarrow$ 0.14% at 2.8m), reflecting site homogenization. For heterogeneous small-scale studies (CV > 5%), practitioners should report both FRI (precision-weighted summary) and $\min\{E_k\}$ (worst-case guarantee) to enable risk-appropriate decision-making. At large homogeneous scales (CV < 1%), this distinction becomes negligible.

3.2 Convergence and Stability Assessment

Numerical convergence pattern: FRI increases from 2.015 (1k) to 2.147 (100k) to 2.149 (2.8m), with changes diminishing at larger scales ($\Delta = 0.132 \rightarrow$

0.002). This pattern suggests convergence, though formal statistical testing is limited by small site count (K=3).

Inter-site homogenization: Coefficient of variation of site E-values collapses from 9.7% (1k) to 0.30% (100k) to 0.16% (2.8m). This sharp reduction indicates that sampling variation dominates at small scales while sites converge to similar E-values at large scales, consistent with sampling theory predictions.

Limitations of bootstrap validation: With only K=3 sites, site-level resampling provides insufficient statistical power for hypothesis testing. Patient-level bootstrap would require access to individual data, violating federated privacy assumptions. Our convergence assessment is thus primarily observational rather than formally tested via inferential statistics, though formal statistical testing faces methodological challenges.

3.3 Computational Performance

Execution Times (2.8m patient dataset):

Operation	Time	Notes
Site-level E-value computation	10s	From MTR bounds
FRI aggregation	<1s	Weighted average
Total	~10s	Practical for real-world use

Scalability: Linear $O(n)$ complexity, consistent with Module 2 results. Memory: ~2-3 GB per site.

3.4 E-value Decomposition Analysis

E-value formula: $E = RR + \sqrt{RR \times (RR - 1)}$

Decomposition into components:

Scale	FRI	RR component	Uncertainty component	Interpretation
1k	1.961	1.25	0.71	High uncertainty from small sample
100k	2.147	1.28	0.87	Moderate uncertainty
2.8m	2.149	1.28	0.87	Stable (converged)

The decomposition reveals that the risk ratio component (treatment effect magnitude) remains stable at approximately 1.28 across all scales, while the uncertainty component increases slightly as bounds tighten with larger samples. This pattern indicates that FRI convergence reflects improvements in statistical precision rather than changes in underlying effect size.

3.5 Illustrative Threshold Application (Unvalidated)

Diabetes Treatment Example (2.8m scale):

Metric	Value	Interpretation
FRI (sample-size weighted)	2.149	Precision-weighted average robustness
$\min\{E_k\}$ (conservative)	2.146	Worst-case site robustness
$\max\{E_k\}$ (optimistic)	2.153	Best-case site robustness

Proposed threshold comparison (Section 2.5):

Threshold Level	Proposed Cutoff	FRI Status	Conservative Status
High-stakes (FDA approval)	>3.0	$2.15 < 3.0$	$2.15 < 3.0$
Moderate (clinical guidelines)	>2.0	$2.15 > 2.0$	$2.15 > 2.0$
Exploratory (research)	>1.5	$2.15 > 1.5$	$2.15 > 1.5$

Interpretation with appropriate caution: Under the **unvalidated** moderate threshold ($FRI > 2.0$), this diabetes treatment effect would be considered “acceptable” for clinical guideline consideration. However, decision-makers should recognize:

1. **Thresholds lack empirical validation:** The 2.0 cutoff is not derived from retrospective analysis of successful clinical decisions or RCT comparisons.
2. **Context matters beyond E-values:** Biological plausibility, study design quality, effect size magnitude, and treatment alternatives all inform decisions. An E-value of 2.15 means a confounder with $RR = 2.15$ (in both treatment and outcome associations) is required to nullify the effect—but whether such a confounder exists depends on domain knowledge.
3. **Conservative alternative:** Risk-averse decision-makers should use $\min\{E_k\} = 2.146$ as the threshold, ensuring even the weakest site’s robustness is respected.
4. **Not a substitute for RCTs:** Observational robustness (E-values) complements but does not replace randomized evidence. Strong E-values suggest unmeasured confounding is unlikely to fully explain effects, but RCTs remain the gold standard for causal inference.

Comparison to known confounders (illustrative): Diabetes treatment confounders typically include disease severity ($RR = 1.8$), medication adherence ($RR = 1.5$), and lifestyle factors ($RR = 1.3$). Since $FRI = 2.15$ exceeds these

plausible values, the treatment effect appears robust to individual unmeasured confounders, though combinations could theoretically reach RR 2.15.

3.6 Communication Efficiency and Privacy Tradeoffs

Table 2: Data Transfer Requirements

Scale	Patients	Centralized	Federated	Reduction
1k	1,130	201 KB	174 bytes	1,156×
100k	235,222	41.9 MB	174 bytes	240,805×
2.8m	2,709,803	482 MB	174 bytes	2.8M×

Per-site transmission (58 bytes): E-value (8), bounds (16), sample size (4), site ID (20), risk ratio (8), metadata (2).

3.6.1 Communication Efficiency $O(1)$ communication complexity:

Federated transmission remains constant at 174 bytes regardless of patient count—a 2,398-fold increase in patients (1k→2.8m) with zero communication increase. This achieves $O(1)$ complexity with respect to patient count, scaling from 1,156× reduction at small scale to 2.8M× at large scale. E-value transmission adds only 8 bytes per site (16% overhead) compared to bounds-only communication, with FRI aggregation performed coordinator-side requiring no additional transmission.

Bandwidth advantage: For networks with K sites and N patients, centralized approaches require $O(K \times N \times p)$ bytes (p = covariate dimension), while federated E-value transmission requires $O(K)$ bytes. This makes billion-scale federated sensitivity analysis feasible on commodity networks.

3.6.2 Covariate Identity Privacy (Primary Design Goal) Methodological privacy preserved:

Sites transmit scalar E-values rather than covariate lists, protecting institutional knowledge about adjustment strategies. In multi-site studies, different sites may use proprietary or strategically valuable variable selections:

- **Example 1 (Psychiatric study):** Site A uses stigmatizing genetic markers (MAO-A variants), Site B socioeconomic factors (income, housing instability), Site C standard compliance measures (medication adherence). Centralized analysis exposes all variable identities to the coordinator and other sites. Federated transmission reveals only E-values, hiding which confounders each site chose.
- **Example 2 (Oncology consortium):** Academic centers may adjust for novel biomarkers under patent consideration, while community hospitals use standard clinical variables. Federated E-values allow participation without disclosing proprietary variable discoveries.

This “covariate identity privacy” has strategic value in competitive research environments and addresses ethical concerns about exposing sensitive variable choices (e.g., socioeconomic proxies, genetic markers).

3.6.3 Information Disclosure Tradeoff (Inherent Design Property)

Quantitative equivalence of E-values and effect sizes: E-values are mathematically equivalent to risk ratios via the monotonic function $E = RR + \sqrt{RR \times (RR - 1)}$. This means:

1. **Invertibility:** Given E, one can solve for RR using $RR = E^2/(2E)$ (approximate) or numerical inversion.
2. **Information content:** Publishing E_k is information-theoretically equivalent to publishing RR_k —both reveal aggregate treatment effect magnitude.
3. **Interpretation:** A coordinator observing $E_k = 2.8$ (high) versus $E_k = 1.6$ (low) can infer the former site has larger effects or adjusted for stronger confounders.

This is a design tradeoff, not a flaw: Meta-analysis fundamentally requires sharing aggregate effect measures. The goal of sensitivity analysis is precisely to quantify and communicate robustness to unmeasured confounding—concealing effect sizes would defeat the purpose. What federated E-values protect is the “how” (covariate identities), not the “what” (aggregate results).

Formal sensitivity analysis: For a sensitivity function $\Delta(E) = \partial RR / \partial E$, the disclosure rate is bounded by $|\Delta| \leq 1$ for $E \in [1.5, 3.0]$, indicating that E-values and RRs carry nearly equivalent information content in this range.

3.6.4 Differential Privacy Extension (Future Work) Roadmap for

formal privacy guarantees: Achieving ϵ -differential privacy requires Laplace noise injection:

$$E'_k = E_k + \text{Lap}\left(0, \frac{\Delta_f}{\epsilon}\right)$$

where: - **Sensitivity Δ_f :** For E-values bounded in $[1, 5]$, global sensitivity $\Delta_f = 4$ (maximum change from adding/removing one patient). - **Privacy budget ϵ :** Standard values range from $\epsilon = 0.1$ (strong privacy) to $\epsilon = 1.0$ (relaxed privacy). - **Utility tradeoff:** With $\epsilon = 1.0$ and $\Delta_f = 4$, Laplace noise has standard deviation $\sigma = 4/1.0 = 4.0$, potentially overwhelming signal ($E = 2.0$).

Calibration challenge: Preliminary analysis suggests $\epsilon = 2.0$ may be required to preserve utility (signal-to-noise ratio > 2), though this requires empirical validation against ground truth confounding scenarios. Adaptive privacy mechanisms (e.g., subsample-and-aggregate, smooth sensitivity) may reduce noise while maintaining guarantees.

3.6.5 Regulatory and Operational Benefits **HIPAA Safe Harbor compliance:** Aggregate statistics without individual identifiers satisfy 45 C.F.R. § 164.514(b), enabling data sharing without full IRB review at each site.

Data Use Agreement simplification: Federated E-values typically do not require Data Use Agreements (DUAs) that govern individual-level data transfers, reducing legal overhead and accelerating multi-site collaboration.

Strategic advantage: By protecting covariate identities while sharing results, federated E-values enable broader participation in consortia where sites have competitive or ethical concerns about exposing methodological choices.

4. DISCUSSION

4.1 Theoretical Implications

Definition 1 characterizes FRI as a mathematically defined aggregation of sensitivity metrics with bounded properties. The boundedness property ($\min < \text{FRI} < \max$) ensures FRI balances optimism and pessimism, though whether this constitutes a formal “guarantee” requires validation against ground truth confounding scenarios.

The **decision-theoretic calibration** (Section 2.5) transforms FRI from a descriptive statistic to a **prescriptive decision tool**, grounded in cost-benefit analysis rather than arbitrary cutoffs.

4.2 Practical Guidelines and Decision Rules

The Federated Robustness Index applies to multi-site observational studies with privacy constraints, heterogeneous treatment effects across institutions, and requirements for robustness quantification without data sharing. The typical workflow proceeds in four steps: compute site-level E-values from local bounds, aggregate using sample-size weighted FRI, compare against decision-theoretic thresholds (Section 2.5), and report findings with explicit threshold interpretation.

For the diabetes study example ($\text{FRI} = 2.15$), appropriate reporting states: “Federated analysis shows moderate robustness ($\text{FRI}=2.15$, exceeding the 2.0 threshold for clinical guidelines), suitable for guideline consideration. An unmeasured confounder would require risk ratio 2.15 in both treatment and outcome associations to explain away the observed treatment effect.” This format provides clinicians with actionable interpretation grounded in decision theory rather than arbitrary cutoffs.

The threshold selection should match decision stakes. High-stakes regulatory decisions (FDA approval) require $\text{FRI} > 3.0$, reflecting stringent evidence standards where Type I errors (approving ineffective treatments) carry severe consequences. Moderate-stakes clinical guidelines appropriately use $\text{FRI} > 2.0$,

balancing evidence quality with clinical utility. Exploratory research contexts accept $FRI > 1.5$, prioritizing hypothesis generation over definitive recommendations.

4.3 Clinical Interpretation Example

Diabetes treatment robustness:

The 2.8m-patient federated analysis yielded $FRI=2.15$, indicating an unmeasured confounder must have $RR\ 2.15$ to nullify the treatment effect.

Comparison to known confounders:

- Disease severity ($RR\sim 1.8$): **Insufficient** to explain effect
- Medication adherence ($RR\sim 1.5$): **Insufficient** to explain effect
- Combined effect ($RR\sim\sqrt{(1.8\times 1.5)}\ 1.64$): **Still insufficient**

Conclusion: Treatment effect is **robust** to plausible unmeasured confounders, supporting clinical guideline inclusion.

4.4 Limitations and Future Directions

Binary outcomes limitation: The current E-value implementation handles binary outcomes exclusively. Extension to continuous outcomes requires modified formulas based on correlation coefficients or standardized mean differences (VanderWeele, 2019). Time-varying treatments necessitate sequential E-value computation (Robins et al., 2000) adapted to federated settings.

Synthetic data validation challenge: Synthea’s complete confounder knowledge creates a fundamental paradox for E-value validation. E-values quantify robustness to unmeasured confounding, but in Synthea all confounders are known by design. We cannot compute “ground truth E-values” because the concept assumes unknown confounders. Real-world validation requires comparing observational study FRI to randomized controlled trial outcomes (where unmeasured confounding is eliminated by design), such as observational hormone therapy studies versus the Women’s Health Initiative RCT. Such validation is beyond our synthetic data scope but represents a critical future direction. Additionally, Synthea simplifies confounding structures compared to real EHR data—real-world multi-site studies often exhibit higher heterogeneity due to geographic variation in treatment practices (Baicker et al., 2013), socioeconomic differences (Chetty et al., 2016), and institutional protocols. Our observed 9.7% inter-site CV at 1k scale likely underestimates real-world heterogeneity.

Site count scalability: Our 3-site validation reflects small hospital consortia but falls short of large networks like FDA Sentinel (18 data partners) or PCORnet (13 clinical research networks). The mathematical characterization (Definition 1) applies to arbitrary K , but empirical heterogeneity patterns and convergence rates require validation across 10+ sites. Future work should assess whether large federated networks exhibit sufficient homogeneity for FRI


convergence or whether regional clustering necessitates hierarchical aggregation strategies.

Privacy architecture strengths and limitations: Section 3.6 establishes that federated E-value transmission achieves two distinct privacy goals with differing guarantees:

1. **Covariate identity privacy (strong):** Sites transmit scalar E-values (174 bytes) rather than covariate lists, hiding which adjustment variables were used. This protects institutional knowledge, addresses ethical concerns about stigmatizing variables, and enables participation in competitive research environments.
2. **Aggregate effect privacy (inherently limited):** E-values are mathematically equivalent to risk ratios via the invertible function $E = RR + \sqrt{RR \times (RR - 1)}$, meaning coordinators can infer effect sizes from E-values. This is a design tradeoff, not a flaw—meta-analysis fundamentally requires sharing aggregate results.

The key insight is that what federated E-values protect is the “**how**” (methodological choices), not the “**what**” (aggregate findings). This aligns with practical privacy needs in multi-site research: sites care more about protecting proprietary variable selections than hiding overall effect estimates that will be published anyway.

Regulatory process validation: While privacy advantages (HIPAA Safe Harbor compliance, Data Use Agreement simplification, covariate identity protection) follow from the federated architecture, claims of IRB timeline improvements lack empirical validation. Future studies should measure actual IRB approval timelines comparing centralized versus federated protocols.



6. CONCLUSIONS

This work proposes the Federated Robustness Index as a sample-size weighted aggregation of site-level E-values for multi-site sensitivity analysis. Validation across three scales (1k-2.8m patients, 3 sites) shows convergence behavior, with FRI increasing from 2.015 to 2.149 as inter-site heterogeneity collapses from 9.7% to 0.16% coefficient of variation.

Methodological contribution: Extending single-site E-value methodology (VanderWeele & Ding, 2017) to federated settings, I characterize FRI’s mathematical properties (Definition 1) and compare aggregation strategies. Sample-size weighting follows meta-analysis principles, giving larger sites proportionally higher weight. The key distinction from prior federated causal inference work (FACE, FLAME, FedCI) lies in aggregating sensitivity metrics rather than

point estimates—a fundamentally different quantity since E-values represent unmeasured confounding thresholds, not effect sizes. The boundedness property ($\min\{E_k\} \leq \text{FRI} \leq \max\{E_k\}$) holds by construction, though FRI provides a precision-weighted summary rather than a formal worst-case guarantee (which would require using $\min\{E_k\}$).

Practical recommendations: For multi-site observational studies, practitioners should report both FRI (precision-weighted summary) and $\min\{E_k\}$ (conservative worst-case), enabling decision-makers to select based on risk tolerance. At small scales with heterogeneity ($\text{CV} > 5\%$), these differ meaningfully (14% gap at 1k scale). At large scales with homogeneity ($\text{CV} < 1\%$), all aggregation strategies converge, making the choice largely irrelevant. When sample sizes vary considerably across sites, sample-size weighting is theoretically justified; when sites are equally sized, simple averaging suffices.

Threshold interpretation caveat: The proposed thresholds (FRI > 3.0 high-stakes, > 2.0 moderate, > 1.5 exploratory) represent exploratory guidelines requiring empirical validation. These are not derived from retrospective analysis of regulatory decisions or RCT comparisons. Decision-makers should treat thresholds as heuristic guidance, not rigid cutoffs, incorporating domain expertise about biological plausibility and study quality beyond statistical metrics. VanderWeele (2019) suggests $E > 2$ indicates “moderate robustness” for single-site studies; our federated thresholds extend this heuristic but inherit its limitations.

Privacy tradeoff architecture: Federated E-value transmission achieves covariate identity privacy (protecting methodological choices) while accepting aggregate effect disclosure (inherent to meta-analysis). Sites transmit 174 bytes versus 482 MB centralized (2.8M $\times$ reduction), hiding which variables were adjusted without concealing overall robustness conclusions. The mathematical equivalence $E \leftrightarrow RR$ means E-values reveal effect magnitudes by necessity—this is a design tradeoff, not a limitation. What matters in practice: sites protect proprietary/sensitive variable selections while sharing results required for scientific transparency. This aligns with real-world privacy needs where institutions care more about protecting “how” (methodology) than “what” (findings). Extensions to formal ϵ -differential privacy ($E'_k = E_k + \text{Lap}(\Delta/\epsilon)$) require calibrating privacy budgets against utility loss (Section 3.6.4), with preliminary analysis suggesting $\epsilon \approx 2.0$ needed to preserve signal-to-noise ratios.

Critical limitations: Synthetic data prevents validation against real unmeasured confounding (see Section 4.4). Our 3-site validation reflects small consortia, not large networks (FDA Sentinel: 18 sites). Threshold calibration lacks empirical grounding—validation priorities include RCT comparison, clinician surveys, and FDA decision analysis. Current implementation handles binary outcomes only.

Honest assessment of contribution: This work does not “prove” federated E-value validity in a formal mathematical sense (Definition 1 characterizes properties of weighted averaging, not robustness preservation). Rather, it proposes

FRI as a principled aggregation method following meta-analysis conventions, characterizes its mathematical properties, and shows convergence on synthetic data. The contribution lies in extending E-value methodology to federated settings with reproducible validation, not in fundamental theoretical breakthroughs. Future work must validate FRI’s practical utility via comparison to RCT ground truth and empirical assessment of proposed thresholds.

5. RELATED WORK

5.1 E-values and Sensitivity Analysis

VanderWeele & Ding (2017) introduced E-values as an intuitive metric for quantifying sensitivity to unmeasured confounding, addressing a longstanding gap in observational causal inference. Ding & VanderWeele (2016) established the mathematical foundations connecting E-values to bias factor formulations. Mathur et al. (2020) developed computational tools (R package ‘EValue’) for practical implementation. Cinelli & Hazlett (2020) proposed sensitivity analysis via `sensemakr`, providing alternative approaches through partial R^2 metrics. Robins et al. (2000) pioneered marginal structural models with sensitivity parameters, laying groundwork for modern sensitivity analysis frameworks.

Our distinction: Existing E-value methodology focuses on single-site analysis. We extend to federated multi-site settings with mathematical characterization of aggregation (Definition 1) and exploratory decision thresholds, enabling privacy-preserving sensitivity analysis without exposing site-specific covariate choices.

5.2 Federated Causal Inference and Aggregation

Recent work addresses federated causal effect estimation but typically assumes unconfoundedness. Li et al. (2022) propose FACE for federated average treatment effects using propensity scores. Zhang et al. (2023) develop FLAME for federated matching. Duan et al. (2020) analyze communication-efficient multi-site regression. Lu et al. (2015) study heterogeneity in meta-analysis aggregation, while DerSimonian & Laird (1986) established random-effects models for combining study-level estimates.

Our distinction: These methods aggregate point estimates or effect sizes, not sensitivity metrics. FRI is the first formal framework for aggregating site-level robustness guarantees in federated settings, providing validity under arbitrary unmeasured confounding.

5.3 Sensitivity Analysis in Multi-Site Studies

Traditional approaches to multi-site sensitivity analysis either pool individual-level data (violating privacy) or use informal qualitative comparison. Rosenbaum (2002) developed sensitivity analysis for matched observational studies

but focused on single-site settings. Cornfield et al. (1959) established early sensitivity analysis principles via “Cornfield’s inequality.” Hosman et al. (2010) proposed sensitivity diagnostics for propensity score matching across subgroups. Carnegie et al. (2016) analyzed sensitivity in multi-site trials with unmeasured confounding.

Our contribution: We provide the first mathematically characterized federated aggregation of sensitivity metrics with boundedness properties (Definition 1), exploratory decision thresholds, and partial covariate privacy (information leakage acknowledged).

5.4 Decision Theory and Threshold Calibration

Classical decision theory (Berger, 1985) provides frameworks for statistical inference under loss functions. Imbens & Manski (2004) apply decision-theoretic principles to partial identification. VanderWeele (2019) discusses thresholds for E-value interpretation but without formal decision-theoretic justification. Our threshold calibration (Section 2.5) extends this work by grounding FRI thresholds in explicit cost-benefit analysis matching clinical stakes.

REFERENCES

1. Rosenbaum, P. R., & Rubin, D. B. (1983). The central role of the propensity score in observational studies. *Biometrika*, 70(1), 41-55.
2. Rosenbaum, P. R. (2002). *Observational studies* (2nd ed.). Springer.
3. Manski, C. F. (2003). *Partial identification of probability distributions*. Springer.
4. Pearl, J. (2009). *Causality: Models, reasoning, and inference* (2nd ed.). Cambridge University Press.
5. VanderWeele, T. J., & Ding, P. (2017). Sensitivity analysis in observational research: introducing the E-value. *Annals of Internal Medicine*, 167(4), 268-274.
6. Ding, P., & VanderWeele, T. J. (2016). Sensitivity analysis without assumptions. *Epidemiology*, 27(3), 368-377.
7. VanderWeele, T. J. (2019). Principles of confounder selection. *European Journal of Epidemiology*, 34(3), 211-219.
8. Mathur, M. B., Ding, P., Riddell, C. A., & VanderWeele, T. J. (2020). Website and R package for computing E-values. *Epidemiology*, 31(2), e26-e28.

9. Cinelli, C., & Hazlett, C. (2020). Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B*, 82(1), 39-67.
10. Cinelli, C., Forney, A., & Pearl, J. (2022). A crash course in good and bad controls. *Sociological Methods & Research*, 51(1), 3-34.
11. Robins, J. M., Hernán, M. Á., & Brumback, B. (2000). Marginal structural models and causal inference in epidemiology. *Epidemiology*, 11(5), 550-560.
12. Cornfield, J., et al. (1959). Smoking and lung cancer: Recent evidence and a discussion of some questions. *Journal of the National Cancer Institute*, 22(1), 173-203.
13. DerSimonian, R., & Laird, N. (1986). Meta-analysis in clinical trials. *Controlled Clinical Trials*, 7(3), 177-188.
14. Lu, G., & Ades, A. E. (2004). Combination of direct and indirect evidence in mixed treatment comparisons. *Statistics in Medicine*, 23(20), 3105-3124.
15. Duan, R., et al. (2020). Learning from electronic health records across multiple sites: A communication-efficient and privacy-preserving distributed algorithm. *Journal of the American Medical Informatics Association*, 27(3), 376-385.
16. Li, S., et al. (2022). Federated causal inference in heterogeneous observational data. *arXiv preprint arXiv:2202.12367*.
17. Zhang, Y., et al. (2023). Privacy-preserving federated causal inference for observational studies. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(12), 14589-14597.
18. Hosman, C. A., Hansen, B. B., & Holland, P. W. (2010). The sensitivity of linear regression coefficients' confidence limits to the omission of a confounder. *Annals of Applied Statistics*, 4(2), 849-870.
19. Carnegie, N. B., Harada, M., & Hill, J. L. (2016). Assessing sensitivity to unmeasured confounding using a simulated potential confounder. *Journal of Research on Educational Effectiveness*, 9(3), 395-420.
20. Berger, J. O. (1985). *Statistical decision theory and Bayesian analysis* (2nd ed.). Springer.
21. Imbens, G. W., & Manski, C. F. (2004). Confidence intervals for partially identified parameters. *Econometrica*, 72(6), 1845-1857.
22. Dwork, C., McSherry, F., Nissim, K., & Smith, A. (2006). Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography Conference* (pp. 265-284). Springer.

23. Baicker, K., et al. (2013). Geographic variation in health care and the problem of measuring racial disparities. *Perspectives in Biology and Medicine*, 56(1), 1-16.
24. Chetty, R., et al. (2016). The association between income and life expectancy in the United States, 2001-2014. *JAMA*, 315(16), 1750-1766.

ETHICS STATEMENT

IRB approval was not required as all data are synthetic (Synthea OMOP CDM v5.4) or publicly available (MIMIC-IV Demo OMOP CDM v5.3). No human subjects were involved.

Data sources: Synthea synthetic patients (1,130 to 2,709,803 across three scales) from AWS Open Data Registry (s3://synthea-omop), and MIMIC-IV Demo (~100 de-identified ICU patients) from PhysioNet (<https://doi.org/10.13026/p1f5-7x35>). Code: <https://github.com/watilde/Harmonia>

End of Manuscript v1.0 (Revised)