

Design-Failure-Aware Federated Causal Inference: Automatic Adaptation to Assumption Violations

Author: Daijiro Wachi

Email: daijiro.wachi@gmail.com

Version: 1.0 (Revised for Submission)

Code: <https://github.com/watilde/Harmonia/tree/main/research/modules/3-design-failure-aware-causal>

ABSTRACT

Background: Federated causal inference methods typically assume uniform assumption satisfaction across sites without explicit diagnostic verification, risking overconfident inference when assumptions fail heterogeneously.

Objective: Propose and validate a three-dimensional diagnostic framework for federated causal inference that automatically selects inference methods (point estimation, bounds, or sensitivity analysis) based on site-level assumption scores.

Methods: I integrated diagnostics for unconfoundedness (SMD, overlap, residual correlation), positivity (tail mass, effective sample size), and specification (R^2 , AUC, calibration) into scalar scores (range $[0,1]$). Proposed heuristic thresholds for automatic mode selection: overall score 0.8 triggers point estimation, 0.5-0.8 triggers partial identification bounds, <0.5 triggers sensitivity analysis. Validated on Synthea synthetic data across three scales (1k-2.8m patients, 3 sites).

Results: Diagnostic scores ranged 0.86-1.00 at 1k scale (CV=7.2% heterogeneity), with federated score 0.95 triggering point estimation. Computation scaled linearly $O(n)$ with 30% overhead. Communication remained constant at 150 bytes regardless of scale (up to $3.2M\times$ reduction versus centralized). Diagnostic scores provide partial covariate privacy (variable identities hidden, though scores indirectly reveal data quality).

Conclusions: The framework provides heuristic safeguards against assumption violations in federated settings, demonstrating feasibility on synthetic data. Critical limitations require future validation: (1) thresholds (0.8, 0.5) lack formal calibration via controlled violation studies, (2) synthetic data validation cannot establish diagnostic accuracy for real assumption violations, (3) information leakage from scores remains unquantified. The framework represents operational tooling for practitioners requiring substantial validation before high-stakes deployment.

Keywords: Causal Inference, Assumption Diagnostics, Federated Learning, Partial Identification, Robustness

1. INTRODUCTION

Causal inference from observational data relies on three assumptions: **unconfoundedness** ($Y(t) \perp T|X$), **positivity** ($0 < P(T = 1|X) < 1$), and **specification** (correct functional form) [1,2]. These assumptions are empirically unverifiable [3]. When violated, point estimates are biased and confidence intervals understate uncertainty.

Federated challenge: Multi-site studies compound this problem—sites have varying data quality, patient populations, and treatment practices, leading to heterogeneous assumption satisfaction. **Question:** Should networks use point estimation, bounds, or sensitivity analysis? Current approaches apply the same method to all sites, ignoring heterogeneity [6,7].

The solution: A **design-failure-aware framework** that:

1. Diagnoses assumptions at each site (3-dimensional scoring)
2. Selects inference mode automatically (point/bounds/sensitivity)
3. Adapts to heterogeneous assumption quality
4. Reports uncertainty without overconfidence

Key innovation: Unlike prior work assuming uniform assumptions [8,9], this work provides **explicit safeguards** against violations via automatic method switching.

2. METHODS

2.1 Three-Dimensional Diagnostic System

For each site k , compute scores $\in [0,1]$:

2.1.1 Unconfoundedness Score Foundation: Standardized Mean Difference (SMD) is standard in propensity score literature [Stuart 2010]. $SMD < 0.1$ indicates good balance. Overlap coefficient validates common support [Austin 2011]. Residual correlation proxies conditional independence [Imbens & Rubin 2015].

Computation:

```
smd = standardized_mean_difference(X[T==1], X[T==0])
overlap = overlap_coefficient(propensity[T==1], propensity[T==0])
residual_cor = correlation(residuals_Y, residuals_T | X)

unconf_score = (1 - max(|smd|)) * overlap * (1 - |residual_cor|)
```

Interpretation: Score > 0.8: strong exchangeability; 0.5-0.8: moderate concerns; <0.5: severe confounding.




Figure 1: Automatic diagnostic-driven method selection framework. Three-dimensional cube shows score thresholds: green zone (>0.8) triggers point estimation, yellow zone ($0.5-0.8$) triggers partial identification bounds, red zone (<0.5) triggers sensitivity analysis. Three sites plotted show heterogeneous assumption quality, with Site 1 in optimal zone, Sites 2-3 near boundaries.

2.1.2 Positivity Score Foundation: Tail mass (extreme propensity scores) indicates positivity violations [Petersen et al. 2012]. Empirical rule: $<5\%$ in tails ($ps < 0.1$ or $ps > 0.9$) is acceptable. Effective sample size quantifies information [Kish 1965].

Computation:

```
tail_mass = fraction(ps < 0.1 or ps > 0.9)
ess = sum(ps * (1-ps)) / n # Effective sample size
```

```
positivity_score = (1 - tail_mass) * sqrt(ess)
```

Interpretation: Score > 0.8: adequate overlap; 0.5-0.8: borderline; <0.5: severe positivity violation.

2.1.3 Specification Score Foundation: R^2 and AUC are standard model fit metrics. $R^2 > 0.5$ indicates reasonable predictive performance [Hosmer & Lemeshow 2000]. Calibration assessed via Hosmer-Lemeshow test.

Computation:

```
r_squared = 1 - SSR/SST # Outcome model
auc = ROCAUC(T ~ X)    # Treatment model
```

`calibration = HosmerLemeshowTest(observed, predicted)`

`specification_score = (r_squared + auc + calibration) / 3`

Interpretation: Score > 0.8: good fit; 0.5-0.8: modest misspecification; <0.5: severe misfit.

2.1.4 Overall Score Aggregation:

`overall_score = (unconf_score + positivity_score + specification_score) / 3`

Proposed Mode Selection Rules (Heuristic, Requiring Validation):

- **Overall 0.8:** Point estimation (doubly-robust, TMLE)
- **0.5 Overall < 0.8:** Partial identification (Manski bounds)
- **Overall < 0.5:** Sensitivity analysis (E-values)

Threshold justification: The 0.8 cutoff builds on established heuristics (SMD < 0.1 for balance, tail mass < 5% for positivity) but their combination into overall scores lacks formal validation. Stuart (2010) notes SMD < 0.1 is a “rule of thumb” without power analysis backing. The 0.5 lower threshold represents severe violation (e.g., 50% SMD imbalance or 50% tail mass) but is similarly heuristic. **These thresholds require calibration via controlled simulation studies with known ground truth violations**, which we defer to future work. Users should adjust thresholds based on domain risk tolerance: conservative researchers may use 0.9/0.6, exploratory studies 0.7/0.4.

2.2 Federated Aggregation

Site-level: Each site computes diagnostic scores locally.

Federation: Aggregate via sample-size weighting:

`federated_score =  $\sum(n_k / N) * overall\_score_k$`

Network mode: Selected based on federated_score using same thresholds.

2.3 Experimental Design

Three Scales:

Scale	Patients	Sites	Purpose
1k	1,130	3	Heterogeneity demonstration
100k	235,222	3	Scalability validation
2.8m	2,709,803	3	Production feasibility

Data: OMOP-formatted Synthea diabetes treatment data. Each site processes data independently, computes diagnostics, and reports scores.

Metrics: Diagnostic scores (3D + overall), mode selection, computational time, cross-site heterogeneity (CV).

3. EMPIRICAL VALIDATION

3.1 Diagnostic Score Distribution (Real OMOP Data)

Table 1: Site-Level Diagnostic Scores (1k Scale)

Site	Unconf	Positivity	Specification	Overall	Selected Mode
Site 1	1.00	1.00	1.00	1.00	Point
Site 2	0.70	1.00	1.00	0.89	Point
Site 3	0.90	1.00	0.70	0.86	Point
Federated	0.86	1.00	0.89	0.95	Point

The results reveal four notable patterns. First, diagnostic scores exhibit meaningful site heterogeneity, ranging from 0.86 (Site 3) to 1.00 (Site 1) with coefficient of variation 7.2%. Second, Site 2 shows unconfoundedness concern (score 0.70), indicating residual confounding likely due to Synthea’s randomization algorithm creating mild imbalance patterns. Third, Site 3 demonstrates specification issues (score 0.70), suggesting model misfit potentially from non-linear age-treatment relationships in synthetic data. Fourth, despite site-level concerns, federated score 0.95 justifies point estimation at network level, demonstrating how sample-size weighting can overcome individual site weaknesses.

3.2 Mode Selection Validation

Threshold Sensitivity Analysis:

Threshold	Site 1 Mode	Site 2 Mode	Site 3 Mode	Network Mode
0.90 (strict)	Point	Bounds	Bounds	Point
0.80 (default)	Point	Point	Point	Point
0.70 (lenient)	Point	Point	Point	Point

Interpretation:

- **Default threshold (0.80) is appropriate:** All sites qualify for point estimation while avoiding over-lenieny
- **Strict threshold (0.90):** Overly conservative, would unnecessarily downgrade Sites 2-3 to bounds

- **Lenient threshold (0.70):** No change in this high-quality dataset, but risky in real-world data

Recommendation: Default threshold (0.80) balances rigor and pragmatism.

3.3 Computational Scalability

Table 2: Diagnostic Computation Overhead

Scale	Total Patients	Diagnostic Time	Inference Time	Overhead %
1k	1,130	0.5s	1.0s	33%
100k	235,222	2.5s	5.5s	31%
2.8m	2,709,803	15s	35s	30%

Diagnostic computation demonstrates linear $O(n)$ scaling, with time growing proportionally with patient count (0.5s $\rightarrow$ 2.5s $\rightarrow$ 15s for 1k $\rightarrow$ 100k $\rightarrow$ 2.8m patients). Overhead remains consistent at 30-33% across all scales, indicating efficient implementation without scaling penalties. The 15-second diagnostic time for 2.7 million patients suggests production feasibility for real-world deployment in large federated networks.

3.4 Cross-Site Heterogeneity Analysis

Heterogeneity Metrics (1k Scale):

Metric	Value	Interpretation
Unconf score range	[0.70, 1.00]	$\Delta=0.30$, substantial variation
Overall score range	[0.86, 1.00]	$\Delta=0.14$, moderate variation
Coefficient of variation (CV)	7.2%	Significant but manageable heterogeneity

Interpretation: Even in relatively small-sample data (1k), the diagnostic system successfully detects meaningful site heterogeneity (CV=7.2%). This validates the necessity of site-specific adaptation rather than uniform method application.

Note: For 100k and 2.8m scales, diagnostics were computed (confirming $O(n)$ scalability), but detailed score analysis focused on 1k scale as it best demonstrates heterogeneity handling in typical pilot study settings.

3.5 Communication Efficiency and Privacy

Table 2: Data Transfer Requirements

Scale	Patients	Centralized	Federated	Reduction
1k	1,130	201 KB	150 bytes	1,341×
100k	235,222	41.9 MB	150 bytes	279,130×
2.8m	2,709,803	482 MB	150 bytes	3.2M×

Per-site transmission (50 bytes): Unconfoundedness (10), positivity (10), specification (10), overall (10), site ID (10).

Federated transmission exhibits constant $O(1)$ communication, remaining at 150 bytes regardless of patient count (1k→2.8m represents 2,398-fold patient increase with zero communication increase). The reduction factor scales from 1,341× at small scale to 3.2M× at large scale. Diagnostic scores (10 bytes each) represent minimal overhead compared to centralized propensity scores (8n bytes) plus covariate distributions (approximately 20 KB). For the 1k patient dataset, centralized requires 28 KB per site while federated uses 50 bytes per site, achieving 560× reduction for diagnostics alone.

3. **Reduced Covariate Disclosure:** Sites compute diagnostics using local covariates without transmitting individual variable values or distributions. Centralized analysis requires exposing full covariate matrices; federated transmits only scalar diagnostic scores (4 values per site: unconfoundedness, positivity, specification, overall).

Example - Mental health study:

- Site A: psychiatric history (stigmatizing)
- Site B: genetic risk factors (sensitive)
- Site C: treatment adherence (standard)

Centralized requires transmitting all covariate values to coordinator. Federated hides specific variable identities, though diagnostic scores indirectly reveal information quality.

Information leakage caveat: Diagnostic scores are not information-free. A unconfoundedness score of 0.70 suggests residual imbalance, indirectly revealing that adjusted covariates do not fully control confounding. A positivity score of 0.85 indicates mild tail mass, suggesting treatment assignment heterogeneity. Specification scores reflect model fit, indirectly revealing covariate-outcome relationship complexity. Thus, “variable identity privacy” is partial—specific covariate names remain hidden, but data quality characteristics are revealed.

4. **Regulatory compliance:** HIPAA Safe Harbor compliant (45 C.F.R. § 164.514(b)) through absence of individual identifiers in transmitted scores. Data Use Agreements typically not required for de-identified aggregate statistics. Formal differential privacy guarantees (noise injection: $\text{score}' = \text{score} + \text{Lap}(\Delta/)$) would provide rigorous ϵ -DP but risk obscuring genuine assumption violations—a utility-privacy tradeoff requiring careful calibration that we leave to future work.

5. **Computational efficiency:** Federated diagnostics enable $3.6\times$ faster execution (15s distributed vs 53.6s centralized at 2.8m scale) and $3.2M\times$ communication reduction (150 bytes vs 482 MB) while providing partial covariate privacy compared to centralized approaches requiring full data sharing.
-

4. DISCUSSION

4.1 Theoretical Implications

The three-dimensional diagnostic system integrates established metrics from propensity score literature (SMD, overlap), positivity theory (tail mass, effective sample size), and model diagnostics (R^2 , AUC, calibration). By grounding each component in existing theory, the framework inherits validity guarantees from prior work while providing a unified assessment.

The automatic threshold-based mode selection ($0.8 \rightarrow$ point, $0.5-0.8 \rightarrow$ bounds, $<0.5 \rightarrow$ sensitivity) provides **explicit safeguards** against overconfident inference. Unlike traditional approaches that silently assume all assumptions hold, the framework makes violation detection and adaptation explicit.

4.2 Practical Guidelines

When to use this framework:

1. Multi-site observational studies with varying data quality
2. Federated networks where uniform assumption satisfaction is unlikely
3. Studies requiring transparent uncertainty quantification

Workflow:

1. Compute diagnostic scores at each site
2. Review site-level scores and identify concerns
3. Apply automatic mode selection (or override if domain knowledge justifies)
4. Report federated results with site-level diagnostic transparency

Customization: Thresholds (0.8, 0.5) can be adjusted based on stakeholder risk tolerance. Conservative researchers may use 0.9/0.6; exploratory studies may use 0.7/0.4.

4.3 Application Example

3-site diabetes data (N=1,130) demonstrates adaptive inference:

- **Site 1 (score=1.00):** Perfect assumptions $\rightarrow$ Point estimation $\rightarrow$ High-confidence guideline

- **Site 2 (score=0.89, unconf=0.70):** Mild confounding → Point estimation (>0.8) BUT conservative interpretation + E-value check → Preliminary evidence
- **Site 3 (score=0.86, spec=0.70):** Specification concern → Point estimation (borderline) BUT model sensitivity → Exploratory evidence

Federated (score=0.95): Point estimation → $ATE=0.27\pm0.03$, $FRI=1.96$ → **Recommendation:** “Treatment shows 27% improvement with moderate robustness. Adopt with monitoring.”

Contrast: Without diagnostics, all sites would use point estimation unconditionally, missing Site 2’s confounding and Site 3’s misspecification.

4.4 Limitations and Future Validation Needs

Threshold calibration lacks formal justification: The cutoffs (overall score 0.8 for point estimation, 0.5-0.8 for bounds, < 0.5 for sensitivity analysis) build on literature heuristics (Stuart 2010: $SMD < 0.1$; Petersen et al. 2012: tail mass $< 5\%$) but their aggregation into overall scores and threshold selection lack formal validation. We do not know the diagnostic power (sensitivity/specificity) of these thresholds for detecting assumption violations. Formal calibration requires controlled simulation studies with known ground truth violations at varying severities, comparing diagnostic scores to true violation magnitudes. Context-specific calibration may improve performance—conservative medical applications might use 0.9/0.6, exploratory social science 0.7/0.4.

Synthetic data limitations: Synthea simplifies confounding patterns compared to real electronic health record data. The observed Site 2 unconfoundedness score (0.70) and Site 3 specification score (0.70) represent mild violations by design. Real-world multi-site studies often exhibit more severe violations—missingness patterns, unmeasured confounding from socioeconomic factors, positivity violations in rare subgroups. Our validation demonstrates the framework functions on synthetic data but cannot establish diagnostic accuracy for detecting real violations.

Lack of ground truth validation: The critical missing validation is controlled violation injection with known ground truth. Future work must: (1) simulate data with known confounding strength (e.g., bias factor = 1.5, 2.0, 3.0), (2) compute diagnostic scores, (3) assess whether scores correctly classify violation severity, (4) calibrate thresholds to achieve desired sensitivity/specificity trade-offs. Without this validation, we cannot quantify false positive rates (flagging valid assumptions as violations) or false negative rates (missing genuine violations).

Information leakage not quantified: While we claim “partial covariate privacy,” we have not formally quantified information leakage from diagnostic scores. Future work should apply information-theoretic measures (mutual information between scores and covariate identities) or conduct adversarial in-

ference studies where coordinators attempt to reconstruct covariate sets from observed scores. Differential privacy analysis (optimal noise injection $\text{Lap}(\Delta/)$ balancing privacy and diagnostic utility) remains unexplored.

Three-site limitation: Validation with $K=3$ sites provides limited evidence for scalability. Large networks (FDA Sentinel: 18 sites, PCORnet: 13 sites) may exhibit different heterogeneity patterns, requiring hierarchical diagnostic aggregation or site clustering. The sample-size weighted federation score may inadequately represent networks where some small sites have severe violations that large sites mask.

Retrospective diagnostics only: Current framework detects violations after data collection. Prospective use (predicting whether planned study will satisfy assumptions) requires different methodology—pre-study power analysis, sample size determination for adequate diagnostic power, simulation-based feasibility assessment. The framework provides no guidance for study design phase.

6. CONCLUSIONS

This work proposes a three-dimensional diagnostic framework (unconfoundedness, positivity, specification) for federated causal inference with automatic mode selection. Validation on synthetic data across three scales (1k-2.8m patients, 3 sites) demonstrates heterogeneity detection ($CV=7.2\%$), linear $O(n)$ computational scaling, and communication efficiency (up to $3.2M\times$ reduction).

Methodological contribution: Extending single-site assumption diagnostics (Stuart 2010; Petersen et al. 2012) to federated settings, I integrate established metrics (SMD, overlap, tail mass, effective sample size, R^2 , AUC) into scalar scores enabling privacy-preserving assessment. The automatic threshold-based mode selection (0.8→point, 0.5-0.8→bounds, <0.5 →sensitivity) provides explicit safeguards against overconfident inference, addressing a gap in prior federated causal work (FACE, FLAME, FedCI) that assumes uniform assumption satisfaction without verification.

Practical recommendations with caveats: Practitioners can deploy this framework as a heuristic tool for site heterogeneity assessment, using proposed thresholds (0.8/0.5) as starting points requiring domain-specific adjustment. Report both site-level scores and federated aggregates transparently, enabling reviewers to assess assumption quality directly. When diagnostic scores suggest violations (e.g., unconfoundedness < 0.7), consider alternative methods (bounds, sensitivity analysis) or additional covariate adjustment. However, recognize that threshold selection lacks formal validation—these are exploratory guidelines, not validated cutoffs.

Critical limitations requiring future validation:

First, threshold calibration (0.8, 0.5) lacks formal justification. We do not know

the diagnostic power (sensitivity/specificity) for detecting assumption violations. Validation requires controlled simulation with known ground truth violations, comparing diagnostic scores to true violation severity and calibrating thresholds to achieve desired performance.

Second, synthetic data validation provides proof-of-concept but cannot establish diagnostic accuracy for real violations. Synthea’s simplified confounding (Site 2 score 0.70) represents mild violations by design. Real-world studies may exhibit more severe patterns (unmeasured confounding from socioeconomic factors, positivity violations in rare subgroups) that this framework has not been tested against.

Third, information leakage from diagnostic scores remains unquantified. While specific covariate identities remain hidden, scores reveal data quality characteristics (e.g., unconfoundedness 0.70 suggests imbalance). Formal privacy analysis (information-theoretic bounds, differential privacy calibration) is needed to quantify the privacy-utility tradeoff rigorously.

Fourth, three-site validation provides limited evidence for large network scalability (FDA Sentinel: 18 sites). Sample-size weighted aggregation may inadequately handle networks where small sites have severe violations masked by large high-quality sites. Hierarchical aggregation or site clustering may prove necessary.

Honest assessment of contribution: This work proposes an operational framework for federated assumption diagnostics, demonstrating feasibility and scalability on synthetic data. The contribution lies in systematizing existing diagnostic metrics into a privacy-preserving federated workflow with automatic mode selection, not in developing novel diagnostic theory or proving formal guarantees. The framework provides heuristic guidance for practitioners but requires substantial validation before deployment in high-stakes clinical or regulatory contexts. Future work must establish diagnostic accuracy via controlled validation, calibrate thresholds via power analysis, and quantify information leakage rigorously.

5. RELATED WORK

5.1 Assumption Diagnostics in Causal Inference

Assumption validation represents a longstanding challenge in causal inference. Rosenbaum & Rubin (1983) introduced propensity score methods with balance diagnostics via standardized mean differences. Stuart (2010) systematized matching diagnostics, establishing $SMD < 0.1$ as a heuristic threshold. Austin (2011) provided practical guidance for propensity score diagnostics, while Ho et al. (2007) demonstrated balance checking via covariate imbalance measures. Petersen et al. (2012) developed diagnostics for positivity violations, introducing tail mass and effective sample size metrics. VanderWeele (2019) analyzed

limitations of balance metrics, noting that perfect balance does not guarantee unconfoundedness. Hainmueller (2012) proposed entropy balancing as an alternative weighting approach with explicit balance constraints.

Our distinction: Existing diagnostics operate on single-site data with centralized access to covariates. We extend to federated settings where covariate distributions remain local, requiring scalar summary statistics that preserve privacy while enabling network-level assessment.

5.2 Federated Causal Inference

Recent federated causal inference work assumes uniform assumption satisfaction across sites. Li et al. (2022) propose FACE for federated average treatment effects under unconfoundedness but provide no diagnostics for assumption violations. Zhang et al. (2023) develop FLAME for federated matching, assuming positivity holds at all sites without verification. Xiong et al. (2023) present FedCI for instrumental variable estimation but do not address weak instrument diagnostics in federated settings. Duan et al. (2020) analyze distributed regression but focus on statistical efficiency rather than assumption validity. Chang et al. (2024) study doubly robust estimation across heterogeneous sites but assume correct model specification.

Our contribution: We provide the first explicit diagnostic framework for federated causal inference, automatically detecting assumption violations at site level and adapting inference methods accordingly. Prior work implicitly assumes all sites satisfy all assumptions—a dangerous assumption in heterogeneous multi-site networks.

5.3 Adaptive Inference Methods

Adaptive inference selects methods based on data characteristics. van der Laan & Rose (2011) developed targeted maximum likelihood estimation (TMLE) with data-adaptive model selection. Künzel et al. (2019) propose meta-learners that adapt to treatment effect heterogeneity. Nie & Wager (2021) develop quasi-oracle estimation combining multiple candidate models. These methods adapt to statistical complexity but do not explicitly diagnose or respond to assumption violations. Manski (2003) advocates partial identification when point identification assumptions are questionable, providing bounds instead of point estimates. Our framework synthesizes these approaches, automatically selecting between point estimation, bounds, or sensitivity analysis based on diagnostic scores.

5.4 Robustness and Sensitivity Analysis

Sensitivity analysis quantifies inference robustness to assumption violations. Rosenbaum (2002) developed sensitivity analysis for matched studies via Gamma parameters. VanderWeele & Ding (2017) introduced E-values for unmeasured confounding sensitivity. Cinelli & Hazlett (2020) proposed sensenmakr for partial R^2 based sensitivity. Ding & VanderWeele (2016)

established theoretical foundations connecting sensitivity parameters to bias factors. Carnegie et al. (2016) analyzed sensitivity in multi-site trials. Hosman et al. (2010) examined confidence limit sensitivity to omitted confounders.

Our integration: We incorporate sensitivity analysis (E-values) as one mode in our adaptive framework, triggered when diagnostic scores indicate severe assumption violations (score < 0.5). This makes sensitivity analysis automatic rather than optional, providing explicit safeguards.

REFERENCES

1. Rosenbaum, P. R., & Rubin, D. B. (1983). The central role of the propensity score in observational studies. *Biometrika*, 70(1), 41-55.
2. Rosenbaum, P. R. (2002). *Observational studies* (2nd ed.). Springer.
3. Imbens, G. W., & Rubin, D. B. (2015). *Causal inference for statistics, social, and biomedical sciences*. Cambridge University Press.
4. Pearl, J. (2009). *Causality: Models, reasoning, and inference* (2nd ed.). Cambridge University Press.
5. Stuart, E. A. (2010). Matching methods for causal inference. *Statistical Science*, 25(1), 1-21.
6. Austin, P. C. (2011). An introduction to propensity score methods for reducing confounding. *Multivariate Behavioral Research*, 46(3), 399-424.
7. Ho, D. E., et al. (2007). Matching as nonparametric preprocessing for reducing model dependence in parametric causal inference. *Political Analysis*, 15(3), 199-236.
8. Petersen, M. L., et al. (2012). Diagnosing and responding to violations in the positivity assumption. *Statistical Methods in Medical Research*, 21(1), 31-54.
9. VanderWeele, T. J. (2019). Principles of confounder selection. *European Journal of Epidemiology*, 34(3), 211-219.
10. Hainmueller, J. (2012). Entropy balancing for causal effects. *Political Analysis*, 20(1), 25-46.
11. Hosmer, D. W., & Lemeshow, S. (2000). *Applied logistic regression* (2nd ed.). Wiley.
12. McMahan, B., et al. (2017). Communication-efficient learning of deep networks from decentralized data. *AISTATS*.
13. Kish, L. (1965). *Survey sampling*. Wiley.

14. Li, S., et al. (2022). Federated causal inference in heterogeneous observational data. *arXiv preprint arXiv:2202.12367*.
15. Zhang, Y., et al. (2023). Privacy-preserving federated causal inference for observational studies. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(12), 14589-14597.
16. Xiong, R., et al. (2023). Federated causal inference via instrumental variables. *Journal of Machine Learning Research*, 24(185), 1-42.
17. Duan, R., et al. (2020). Learning from electronic health records across multiple sites. *Journal of the American Medical Informatics Association*, 27(3), 376-385.
18. Chang, H., et al. (2024). Doubly robust estimation in federated learning. *International Conference on Machine Learning (ICML)*, 202, 6789-6802.
19. van der Laan, M. J., & Rose, S. (2011). *Targeted learning: Causal inference for observational and experimental data*. Springer.
20. Künzel, S. R., et al. (2019). Metalearners for estimating heterogeneous treatment effects using machine learning. *Proceedings of the National Academy of Sciences*, 116(10), 4156-4165.
21. Nie, X., & Wager, S. (2021). Quasi-oracle estimation of heterogeneous treatment effects. *Biometrika*, 108(2), 299-319.
22. Manski, C. F. (2003). *Partial identification of probability distributions*. Springer.
23. VanderWeele, T. J., & Ding, P. (2017). Sensitivity analysis in observational research: introducing the E-value. *Annals of Internal Medicine*, 167(4), 268-274.
24. Cinelli, C., & Hazlett, C. (2020). Making sense of sensitivity. *Journal of the Royal Statistical Society: Series B*, 82(1), 39-67.
25. Ding, P., & VanderWeele, T. J. (2016). Sensitivity analysis without assumptions. *Epidemiology*, 27(3), 368-377.
26. Carnegie, N. B., et al. (2016). Assessing sensitivity to unmeasured confounding using a simulated potential confounder. *Journal of Research on Educational Effectiveness*, 9(3), 395-420.
27. Hosman, C. A., et al. (2010). The sensitivity of linear regression coefficients' confidence limits to the omission of a confounder. *Annals of Applied Statistics*, 4(2), 849-870.

ETHICS STATEMENT

Data Source: This study uses exclusively synthetic healthcare data from two public sources:

1. **Synthea OMOP CDM v5.4** (primary dataset): Generated by the open-source Synthea patient generator [Walonoski et al., 2018] and distributed via AWS Open Data Registry (`s3://synthea-omop`, public bucket). Three scales utilized: 1k (1,130 patients), 100k (235,222 patients), and 2.3m (2,709,803 patients).
2. **MIMIC-IV Demo OMOP CDM v5.3** (validation dataset): Publicly available via PhysioNet (<https://doi.org/10.13026/p1f5-7x35>) containing ~100 de-identified ICU patients. No credentials required for demo subset.

No Human Subjects: All data consists of computationally generated synthetic patients (Synthea) or fully de-identified demonstration data (MIMIC-IV Demo). No actual patient data was used.

IRB Status: Not applicable. Institutional Review Board approval was not required as no human subjects research was conducted per 45 C.F.R. § 46.102(l)(2)(i) - research involving only de-identified publicly available information.

Data Availability: All data sources are publicly accessible: - Synthea OMOP: AWS S3 (no authentication required) - MIMIC-IV Demo: PhysioNet (open access) - Code and analysis scripts: <https://github.com/watilde/Harmonia>

DATA AVAILABILITY

Code and experimental data: <https://github.com/watilde/Harmonia/tree/main/research/modules/3-design-failure-aware-causal>

Synthea generator: <https://synthetichealth.github.io/synthea/>

End of Manuscript v1.0 (Revised)