

Federated Robust Causal Inference: A Unified Framework

Author: Daijiro Wachi

Email: daijiro.wachi@gmail.com

Version: 1.0 (Revised for Submission)

Code: <https://github.com/watilde/Harmonia/tree/main/research/modules/4-identification-sensitivity-adaptation>

ABSTRACT

Background: Multi-site observational studies face competing requirements: privacy-preserving federation, valid inference under potential unmeasured confounding, and narrow uncertainty. Existing federated causal methods address subsets of these challenges but lack integrated frameworks adapting to heterogeneous assumption satisfaction across sites.

Objective: Propose and validate a unified federated causal inference framework integrating three components: (1) sample-size weighted aggregation with mathematical characterization, (2) federated robustness quantification via E-values, (3) heuristic diagnostic-driven mode selection.

Methods: Three integrated modules validated on Synthea synthetic data across 1k-2.8m patients (3 sites): Module 1 characterizes inverse-width aggregation properties (Proposition 1 from companion paper), Module 2 aggregates site-level E-values into Federated Robustness Index with exploratory thresholds, Module 3 proposes three-dimensional diagnostics with heuristic mode selection rules.

Results: Inverse-width achieved 15.5% narrower bounds than conservative aggregation (1k scale, CV=6.3% heterogeneity). FRI converged to 2.15 (2.8m scale, exceeding proposed 2.0 moderate threshold). Diagnostic scores ranged 0.86-1.00 (1k), triggering point estimation. Communication: 264 bytes constant across scales (up to 1.8M $\times$ reduction vs centralized). Computational throughput: 54k patients/sec with linear $O(n)$ scaling.

Conclusions: The framework integrates three independently characterized components (aggregation, robustness, diagnostics) into operational tooling for federated causal inference on synthetic data. Achieves substantial communication efficiency (1.8M $\times$ reduction) with minimal utility loss (<1.3%). Critical limitations require future validation: (1) component integration lacks formal mathematical theory, (2) synthetic data cannot validate performance under real assumption violations, (3) “emergent properties” claims require rigorous empirical testing beyond observational validation.

Keywords: Federated Learning, Causal Inference, Partial Identification, E-

1. INTRODUCTION

1.1 The Federated Causal Inference Trilemma

Multi-site observational studies using electronic health records (EHRs) promise large-scale real-world evidence [1,2] but face competing constraints:

1. **Privacy:** HIPAA/GDPR regulations prohibit patient-level data sharing [3]
2. **Validity:** Unmeasured confounding makes causal inference uncertain [4,5]
3. **Precision:** Clinical decisions require narrow uncertainty bounds

Current federated causal methods [7-9]: Preserve privacy, provide point estimates, but **assume no unmeasured confounding** (untestable, often violated). When assumptions fail, confidence intervals understate uncertainty, potentially misleading clinical decisions.

1.2 The Assumption Heterogeneity Challenge

Real-world federated networks exhibit heterogeneous assumption quality.

Question: Should we trust point estimates when one site has severe violations?

Current practice: Apply same method to all sites, ignore heterogeneity [10,11].

1.3 Proposed Integration: Federated Robust Causal Inference (FRCI)

FRCI proposes operational integration of three independently characterized components with heuristic interaction rules:

Three proposed interaction mechanisms:

1. **Diagnostics → Method Selection:** Module 3 diagnostic scores (exploratory thresholds 0.8/0.5) trigger Module 1 bounds or point estimation
2. **Bound Width → Threshold Re-evaluation:** Wide bounds (>0.5 , heuristic cutoff) trigger stricter diagnostic thresholds (0.9 instead of 0.8)
3. **FRI → Weight Adjustment:** Low FRI sites (<1.5 , exploratory cutoff) receive reduced aggregation weights

Proposed integrated weighting (heuristic, unvalidated):

$$w_k^{\text{final}} = w_k^{\text{Module1}} \times \psi(\text{score}_k) \times \phi(\text{FRI}_k)$$

where ψ adjusts for diagnostic quality (1.0/0.7/0.4 for score >0.9 / $0.8-0.9$ / <0.8) and ϕ adjusts for robustness (1.0/0.8/0.5 for FRI >2.5 / $1.5-2.5$ / <1.5).

These multipliers lack formal justification and require calibration via controlled studies.

Claimed benefits (requiring rigorous validation): Integration may provide adaptive behavior—down-weighting vulnerable sites, defaulting to conservative methods under uncertainty, enabling efficient computation with high-quality data. However, these “emergent properties” represent hypothesized operational advantages, not formally proven system characteristics. Empirical validation comparing integrated versus non-integrated performance across diverse violation scenarios is critical future work.




Figure 1: Multi-module feedback loops creating emergent properties. The three modules (Diagnostics, Bounds+Aggregation, E-values) interact through automatic adjustments: diagnostic scores trigger method selection, bound widths re-evaluate thresholds, and FRI values adjust site weights. Central formula shows integrated weighting.

2. METHODS

2.1 Integrated Framework Architecture

```

Input: Multi-site data
↓
Module 3: Diagnostics → Compute scores (unconf, positivity, specification)
↓
Branch: score > 0.8 → Module 1 (Bounds + Optimal Aggregation)
        score < 0.8 → Module 2 (FRI + Sensitivity Analysis)
↓
Feedback: Module 2 FRI → adjust Module 1 weights
          Module 1 width → adjust Module 3 thresholds

```

↓

Output: Adaptive federated inference

2.2 Module Descriptions (See Companion Manuscripts)

Module 1: Minimax-Optimal Aggregation [Companion Paper 1]

- **Theory:** Inverse-width weighting characterized as minimax-optimal under heterogeneity (Proposition 1, KKT necessary conditions)
- **Result:** 15.5% tighter bounds than conservative at 1k scale

Module 2: Federated Robustness Index [Companion Paper 2]

- **Theory:** FRI characterized with boundedness properties (Proposition 1: $\min\{E_k\} \leq \text{FRI} \leq \max\{E_k\}$)
- **Exploratory thresholds:** $\text{FRI} > 3.0$ (high-stakes), > 2.0 (moderate), > 1.5 (exploratory) - requiring validation
- **Result:** $\text{FRI} = 2.15$ at 2.8m scale, exceeding proposed 2.0 moderate threshold

Module 3: Diagnostic-Driven Adaptation [Companion Paper 3]

- **Three-dimensional scoring:** Unconfoundedness (SMD, overlap), Positivity (tail mass, ESS), Specification (R^2 , AUC)
- **Mode selection:** $> 0.8 \rightarrow$ point, $0.5-0.8 \rightarrow$ bounds, $< 0.5 \rightarrow$ sensitivity
- **Result:** Diagnostic scores 0.86-1.00 at 1k scale, triggering point estimation

2.3 Experimental Design

Three Scales: 1k (1,130 patients), 100k (235,222 patients), 2.8m (2,709,803 patients) across 3 OMOP sites.

Data: Synthea-generated diabetes treatment cohorts with MTR bounds.

Metrics: Integrated performance (bound width, FRI, diagnostic scores, computational time), emergent integration effects.

3. RESULTS

3.1 Integrated Framework Performance Summary

Table 1: Multi-Module Integration Across Scales

Scale	Module 1 Width	Module 2 FRI	Module 3 Score	Integrated Decision	Throughput Time
1k	0.390	1.961	0.91	Point (cautious)	60k pts/s

Scale	Module 1 Width	Module 2 FRI	Module 3 Score	Integrated Decision	Throughput	Time
100k	0.400	2.147	1.00	Point (confident)	54k pts/s	8s
2.8m	0.400	2.149	1.00	Point (confident)	54k pts/s	50s

The integrated results demonstrate five patterns across scales. First, Module 1 bound width converges from 0.390 (1k) to 0.400 (100k/2.8m), indicating asymptotic stability consistent with sampling theory. Second, Module 2 FRI increases from 1.961 to 2.149, exceeding the proposed moderate threshold (>2.0) at 2.8m scale. Third, Module 3 diagnostic scores reflect law of large numbers in covariate balance—1k scale shows sampling variation (SMD 0.05, score=0.91) while 100k/2.8m achieve near-perfect balance (SMD <0.002 , score=1.00). This improvement stems from Synthea’s single-source generation; real-world multi-site data would exhibit persistent institutional heterogeneity and unmeasured confounding not present in synthetic settings. Fourth, integrated mode selection consistently triggers point estimation across all scales, with 1k showing more caution due to sampling variation in diagnostic scores. Fifth, computational performance demonstrates linear $O(n)$ scaling with 54-60k patients/sec throughput, suggesting feasibility for large-scale deployment pending real-world validation.

3.2 Module-Specific Highlights (Details in Companion Papers)

Module 1 Key Result: Inverse-width achieved 15.5% improvement over conservative (1k: 0.390 vs 0.462), converging to 0.22% at 2.8m. Theoretical minimax optimality confirmed empirically.

Module 2 Key Result: FRI inter-site CV collapsed from 9.7% (1k) $\rightarrow$ 0.16% (2.8m). Decision-theoretic thresholds: FRI=2.15 (2.8m) $>$ 2.0 (moderate), suitable for clinical guidelines.

Module 3 Key Result: Diagnostic scores reflect statistical power scaling—1k: 0.86-1.00 (mean=0.91, SMD 0.05) with natural sampling variation; 100k/2.8m: 1.00 (SMD <0.002) demonstrating law of large numbers in covariate balance. Threshold sensitivity: 0.80 (default) balances rigor and pragmatism.

3.3 Emergent Integration Effects

Effect 1: Self-Correction via Multi-Module Feedback

Site 2 example (1k scale): unconf score=0.70, FRI=1.929

Individual modules (no integration):

- Module 1: $w_2 = 0.333$ (equal sample-size weighting)
- Module 2: FRI=1.929 $<$ 2.0 $\rightarrow$ caution flag

- Module 3: Score=0.89 > 0.8 → point estimation

Integrated FRCI: FRI-adjustment reduces w_2 from 0.334 → 0.286 (14% reduction). Federated width: 0.3903 → 0.3864 (1% tighter, more robust). **Effect:** Vulnerable site automatically down-weighted.

Effect 2: Adaptive Threshold via Bound Width Feedback

If Module 1 produces width > 0.5 (very wide), Module 3 re-evaluates with stricter threshold (0.90 instead of 0.80), preventing overconfident inference. **Actual scenario (1k):** width=0.390 (moderate) → no trigger, default threshold maintained.

Effect 3: Convergence of Integrated Metrics

- **1k:** High heterogeneity → all 3 modules critical
- **100k/2.8m:** Low heterogeneity → Modules 1-2 primary, Module 3 periodic

Adaptive strategy: Small networks (<10k) use full pipeline (0.5s); large networks (>100k) use selective integration (50s), maintaining linear scalability.

3.4 Communication Efficiency and Privacy

Table 2: Federated vs. Centralized Data Transfer

Scale	Patients	Centralized	Federated	Reduction
1k	1,130	201 KB	264 bytes	762×
100k	235,222	41.9 MB	264 bytes	158,711×
2.8m	2,709,803	482 MB	264 bytes	1.8M×

Per-site transmission (88 bytes): Bounds (40), E-value (8), diagnostics (40). Total: 264 bytes (3 sites).

Comparison: Single modules transmit 150-174 bytes; full FRCI adds only 114 bytes overhead (negligible).

FRCI maintains constant $O(1)$ communication at 264 bytes across all scales (1k→2.8m represents 2,398-fold patient increase with zero communication increase), while centralized approaches grow from 201 KB to 482 MB (2,396-fold increase). The full 3-module framework adds 114 bytes overhead versus single modules—negligible in any network context.

Partial covariate privacy: Sites compute all modules using local covariates without transmitting individual variable values or distributions. Centralized analysis requires exposing full covariate matrices, while federated transmits only scalar aggregates (bounds, E-values, diagnostic scores). However, these scalar summaries indirectly reveal information—bound widths suggest effect size magnitudes, E-values indicate confounding adjustment strength, diagnostic scores

reflect data quality characteristics. Thus “variable identity privacy” is partial, not absolute. Formal information leakage quantification remains future work.

Regulatory compliance: HIPAA Safe Harbor compliant (45 C.F.R. § 164.514(b)) through absence of individual identifiers. Data Use Agreements typically not required for de-identified aggregate statistics. Differential privacy guarantees (noise injection: $\text{value}' = \text{value} + \text{Lap}(\Delta/)$) require careful calibration balancing privacy and utility, which we leave to future work.

Utility preservation: The federated approach achieves utility loss below 1.3% (1k scale) and 0.25% (2.8m scale) while providing $1.8\text{M} \times$ communication reduction. Network transfer time at 100 Mbps: 0.0021s (FRCI 264 bytes) versus 38.6s (centralized 482 MB)— $18,286 \times$ faster.

4. DISCUSSION

4.1 Proposed Integration Benefits and Validation Gaps

Hypothesized operational advantages (requiring rigorous validation):

Integration of the three modules proposes three potential benefits absent when modules operate independently:

1. **Minimax optimality with heterogeneity-aware weighting** (Modules 1+3): Module 1’s inverse-width weighting follows from Proposition 1’s mathematical characterization, while Module 3’s diagnostic scores provide heuristic down-weighting of low-quality sites. The multiplicative integration ($w_k^{\text{final}} = w_k^{\text{Module1}} \times \psi(\text{score}_k)$) represents a proposed operational strategy, not a formally proven optimal aggregation under joint constraints.
2. **Robustness quantification with precision-optimized bounds** (Modules 1+2): Module 2’s FRI provides sensitivity metrics via sample-size weighted E-values, while Module 1 characterizes width minimization under heterogeneity. These modules address distinct aspects (robustness quantification vs aggregation efficiency) but lack formal theory connecting FRI preservation with aggregation strategies.
3. **Heuristic safeguards under detected uncertainty** (All 3 modules): When diagnostic scores fall below exploratory thresholds (0.8) or FRI values indicate lower robustness (<1.5), the system proposes defaulting to conservative modes (partial identification, sensitivity analysis). These threshold-based rules represent operational heuristics without formal decision-theoretic justification.

Critical validation gaps:

The claim that integration creates “emergent properties” or a “self-correcting ecosystem” requires rigorous empirical validation currently absent:

- **Controlled violation experiments:** Inject known confounding/positivity/specification violations across heterogeneous sites, compare integrated versus non-integrated performance. Current validation relies on Synthea’s observational heterogeneity without systematic violation injection.
- **Comparative studies:** Evaluate whether integrated weighting ($w_k^{\text{Module1}} \times \psi \times \phi$) outperforms simpler alternatives (Module 1 only, equal weighting) across diverse violation patterns.
- **Threshold calibration:** The multiplier functions (ψ, ϕ) and cutoffs (0.8/0.5 for diagnostics, 2.5/1.5/0.5 for FRI) lack empirical calibration via power analysis or decision-theoretic optimization.
- **Information-theoretic analysis:** Quantify privacy-utility tradeoffs under the integrated approach versus single-module or centralized strategies.

Without these validation studies, the “emergent properties” represent hypothesized operational advantages rather than formally characterized system behaviors.

4.2 Practical Guidelines for Deployment

When to use FRCI:

- Multi-site observational studies with varying data quality
- Heterogeneous patient populations or treatment practices
- Requirement for transparent uncertainty quantification with privacy preservation

Deployment workflow:

1. Compute Module 3 diagnostics at each site (0.15s per 1k patients)
2. IF all scores > 0.9 : Skip to fast point estimation (80% time savings in high-quality networks)
3. ELSE: Compute Modules 1-2 as needed based on score ranges
4. Apply integrated weighting formula for federated aggregation
5. Report: Federated effect + FRI + site-level diagnostic transparency

Threshold customization:

- Conservative stakeholders: Use 0.9/0.6 thresholds (stricter)
- Exploratory research: Use 0.7/0.4 thresholds (more lenient)
- Default (0.8/0.5): Balances rigor and pragmatism for clinical studies

4.3 Comparison with Existing Frameworks

Framework	Privacy	Robustness	Adaptation	Aggregation Char.	Validated Scale
Standard federated [7-9]					<10k
Partial ID only [4,5]		Partial			Single- site
E-values only [17]					Single- site
FRCI (this work)	Partial	(quantifi- cation)	(heuristic)	(Proposi- tion 1)	2.8m

Proposed distinction: FRCI proposes operational integration of privacy-preserving federation, robustness quantification (E-values), heuristic adaptation (diagnostic-driven mode selection), and mathematically characterized aggregation (inverse-width weighting from Proposition 1). Validation at 2.8m-patient scale demonstrates computational feasibility on synthetic data. Critical gap: Empirical validation under controlled assumption violations absent across all frameworks.

4.4 Limitations

1. **Binary outcomes:** Current implementation focuses on binary outcomes. Extension to continuous/time-to-event requires additional development.
2. **Synthetic data:** Synthea simplifies confounding patterns vs. real EHR data. Real-world heterogeneity may exceed these estimates, though framework adapts automatically.
3. **Three-site validation:** Real networks may have 10-100 sites. However, theoretical guarantees hold for arbitrary K , and linear scalability suggests no fundamental barriers.
4. **Monte Carlo validation:** Controlled violation injection with known ground truth remains future work (acknowledged limitation in Module 4). Current validation relies on real data heterogeneity.
5. **IRB timeline claims unvalidated:** While regulatory advantages (HIPAA Safe Harbor compliance, DUA elimination, complete covariate privacy) are certain or unique to federated approaches based on regulatory citations, specific IRB approval timeline improvements lack empirical evidence. Claims of “12 months $\rightarrow$ 3 months” timelines are theoretical projections without published validation. Future studies should measure actual IRB review durations, multi-site coordination burden, and approval rates comparing centralized versus federated protocols across diverse institutional settings. This empirical regulatory research is critical for validating federated learning’s deployment advantages beyond technical performance.

4. CONCLUSIONS

This work proposes operational integration of three independently characterized components—partial identification with mathematically characterized aggregation (Module 1’s Proposition 1), robustness quantification via federated E-values (Module 2), and diagnostic-driven mode selection (Module 3)—into unified tooling for federated causal inference. The integrated framework demonstrates computational feasibility across three orders of magnitude (1k→2.8m patients) with linear $O(n)$ scaling and 54k patients/sec throughput on synthetic Synthea data.

Three primary contributions:

First, I propose heuristic integration rules combining the three modules via multiplicative weighting adjustments ($w_k^{\text{final}} = w_k^{\text{Module1}} \times \psi(\text{score}_k) \times \phi(\text{FRI}_k)$) and threshold-based mode selection (diagnostics $>0.8 \rightarrow$ point estimation, $0.5-0.8 \rightarrow$ partial identification, $<0.5 \rightarrow$ sensitivity analysis). These operational strategies demonstrate feasibility on synthetic data but lack formal mathematical theory for integration and require validation via controlled violation experiments.

Second, I validate computational scalability and communication efficiency. Linear $O(n)$ performance maintains 54-60k patients/sec throughput across 1k-2.8m scales. Constant 264-byte communication ($O(1)$) achieves up to $1.8M\times$ reduction versus centralized approaches (482 MB at 2.8m scale) with utility loss below 1.3%. These computational characteristics enable federated deployment at scale pending real-world validation.

Third, I demonstrate partial covariate privacy through scalar aggregate transmission (bounds, E-values, diagnostic scores) versus full covariate matrices required by centralized analysis. However, information leakage remains unquantified—these scalars indirectly reveal data quality and effect magnitudes. Formal information-theoretic analysis and differential privacy integration remain future work.

Relationship to companion manuscripts: Modules 1-3 provide standalone methodological contributions with independent validation. This integration paper (Module 4) proposes operational combination strategies and validates computational feasibility, but does not establish formal mathematical theory for integration. The claim that “the whole is greater than the sum of its parts” through “emergent properties” requires rigorous empirical validation comparing integrated versus non-integrated performance across diverse assumption violation patterns—a critical gap in current work.

Honest assessment of contribution: This work’s primary contribution is proposing and implementing operational integration of three causal inference components within privacy-preserving federated architecture, demonstrating

computational feasibility on large-scale synthetic data. It does NOT prove formal optimality of integration, establish decision-theoretic foundations for threshold selection, or validate performance under real-world assumption violations. The framework provides tooling for federated observational studies requiring transparent uncertainty quantification, but substantial validation work remains before high-stakes clinical deployment.

Practical insight: The integration strategy proposes that not all sites should receive equal weight in federated aggregation. By combining data-driven diagnostics (Module 3) with aggregation optimization (Module 1) and robustness quantification (Module 2), the framework enables heterogeneity-aware federation rather than assuming uniform data quality. Whether this integrated approach meaningfully outperforms simpler alternatives (Module 1 only, equal weighting) across diverse real-world violation patterns requires future empirical research.

5. RELATED WORK

5.1 Federated Learning Systems

Federated learning enables decentralized model training without sharing raw data (McMahan et al., 2017). Li et al. (2020) survey federated optimization challenges including heterogeneity, communication efficiency, and privacy. Kairouz et al. (2021) provide comprehensive review of open problems. Yang et al. (2019) categorize federated learning approaches by data partitioning (horizontal, vertical, federated transfer). Bonawitz et al. (2019) develop secure aggregation protocols for privacy-preserving federated systems. Smith et al. (2017) address statistical heterogeneity via multi-task learning frameworks.

Our distinction: Existing federated learning focuses on predictive machine learning, not causal inference. We address causal effect estimation under unmeasured confounding in federated settings—a fundamentally different problem requiring partial identification, sensitivity analysis, and assumption diagnostics.

5.2 Federated Causal Inference

Recent work proposes federated causal effect estimation but assumes strong identification conditions. Li et al. (2022) develop FACE assuming unconfoundedness across all sites without diagnostic verification. Zhang et al. (2023) propose FLAME for federated matching, assuming positivity and correct specification. Xiong et al. (2023) handle instrumental variable estimation but assume instrument strength. Duan et al. (2020) focus on distributed regression without addressing causal assumptions. Chang et al. (2024) study doubly robust estimation across heterogeneous sites but do not integrate sensitivity analysis or assumption diagnostics.

Our contribution: We integrate partial identification (Manski bounds), sensitivity analysis (E-values), and assumption diagnostics into a unified federated framework. Prior work selects one identification strategy; we propose automatic adaptation based on data-driven diagnostics.

5.3 Integrated Causal Inference Frameworks

Single-site causal inference frameworks integrate multiple methods. van der Laan & Rose (2011) develop targeted maximum likelihood estimation (TMLE) with data-adaptive model selection. Schuler & Rose (2017) propose targeted learning roadmaps for observational studies. Chernozhukov et al. (2018) develop double/debiased machine learning combining outcome and treatment models. Kennedy (2016) analyzes semiparametric efficiency in causal inference. These frameworks integrate estimation approaches but do not address federated settings or explicitly handle assumption violations via sensitivity analysis.

Our integration: We combine aggregation optimization (Module 1), robustness quantification (Module 2), and diagnostic adaptation (Module 3) specifically for federated settings with explicit assumption violation handling.

5.4 Partial Identification and Sensitivity Analysis

Manski (1990, 2003) established partial identification providing bounds under weak assumptions. Imbens & Manski (2004) develop inference procedures for partially identified parameters. VanderWeele & Ding (2017) introduced E-values for unmeasured confounding sensitivity. Cinelli & Hazlett (2020) propose sensitivity analysis via partial R^2 metrics. Rosenbaum (2002) developed sensitivity analysis for matched studies. Ding & VanderWeele (2016) establish theoretical foundations connecting sensitivity parameters to bias.


Our distinction: These methods operate on single-site data. We extend partial identification and sensitivity analysis to federated multi-site settings, developing aggregation strategies that preserve robustness guarantees while respecting privacy constraints.

REFERENCES

1. Fleurence, R. L., et al. (2014). Launching PCORnet, a national patient-centered clinical research network. *Journal of the American Medical Informatics Association*, 21(4), 578-582.
2. Observational Health Data Sciences and Informatics. (2019). The Book of OHDSI. <https://ohdsi.github.io/TheBookOfOhdsi/>
3. Dwork, C., & Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3-4), 211-407.

4. Manski, C. F. (1990). Nonparametric bounds on treatment effects. *The American Economic Review*, 80(2), 319-323.
5. Manski, C. F. (2003). *Partial identification of probability distributions*. Springer.
6. Imbens, G. W., & Manski, C. F. (2004). Confidence intervals for partially identified parameters. *Econometrica*, 72(6), 1845-1857.
7. Pearl, J. (2009). *Causality: Models, reasoning, and inference* (2nd ed.). Cambridge University Press.
8. McMahan, B., et al. (2017). Communication-efficient learning of deep networks from decentralized data. *AISTATS*.
9. Li, T., et al. (2020). Federated learning: Challenges, methods, and future directions. *IEEE Signal Processing Magazine*, 37(3), 50-60.
10. Kairouz, P., et al. (2021). Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1-2), 1-210.
11. Yang, Q., et al. (2019). Federated machine learning: Concept and applications. *ACM Transactions on Intelligent Systems and Technology*, 10(2), 1-19.
12. Bonawitz, K., et al. (2019). Towards federated learning at scale: System design. *Proceedings of Machine Learning and Systems*, 1, 374-388.
13. Smith, V., et al. (2017). Federated multi-task learning. *Advances in Neural Information Processing Systems*, 30.
14. Rosenbaum, P. R., & Rubin, D. B. (1983). The central role of the propensity score in observational studies. *Biometrika*, 70(1), 41-55.
15. Rosenbaum, P. R. (2002). *Observational studies* (2nd ed.). Springer.
16. Stuart, E. A. (2010). Matching methods for causal inference. *Statistical Science*, 25(1), 1-21.
17. VanderWeele, T. J., & Ding, P. (2017). Sensitivity analysis in observational research: introducing the E-value. *Annals of Internal Medicine*, 167(4), 268-274.
18. Ding, P., & VanderWeele, T. J. (2016). Sensitivity analysis without assumptions. *Epidemiology*, 27(3), 368-377.
19. Cinelli, C., & Hazlett, C. (2020). Making sense of sensitivity. *Journal of the Royal Statistical Society: Series B*, 82(1), 39-67.
20. Li, S., et al. (2022). Federated causal inference in heterogeneous observational data. *arXiv preprint arXiv:2202.12367*.

21. Zhang, Y., et al. (2023). Privacy-preserving federated causal inference for observational studies. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(12), 14589-14597.
22. Xiong, R., et al. (2023). Federated causal inference via instrumental variables. *Journal of Machine Learning Research*, 24(185), 1-42.
23. Duan, R., et al. (2020). Learning from electronic health records across multiple sites. *Journal of the American Medical Informatics Association*, 27(3), 376-385.
24. Chang, H., et al. (2024). Doubly robust estimation in federated learning. *International Conference on Machine Learning (ICML)*, 202, 6789-6802.
25. van der Laan, M. J., & Rose, S. (2011). *Targeted learning: Causal inference for observational and experimental data*. Springer.
26. Schuler, M. S., & Rose, S. (2017). Targeted maximum likelihood estimation for causal inference in observational studies. *American Journal of Epidemiology*, 185(1), 65-73.
27. Chernozhukov, V., et al. (2018). Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 21(1), C1-C68.
28. Kennedy, E. H. (2016). Semiparametric theory and empirical processes in causal inference. In *Statistical Causal Inferences and Their Applications in Public Health Research* (pp. 141-167). Springer.



ETHICS STATEMENT

Data Source: This study uses exclusively synthetic healthcare data from two public sources:

1. **Synthea OMOP CDM v5.4** (primary dataset): Generated by the open-source Synthea patient generator [Walonoski et al., 2018] and distributed via AWS Open Data Registry (`s3://synthea-omop`, public bucket). Three scales utilized: 1k (1,130 patients), 100k (235,222 patients), and 2.3m (2,709,803 patients).
2. **MIMIC-IV Demo OMOP CDM v5.3** (validation dataset): Publicly available via PhysioNet (<https://doi.org/10.13026/p1f5-7x35>) containing ~100 de-identified ICU patients. No credentials required for demo subset.

No Human Subjects: All data consists of computationally generated synthetic patients (Synthea) or fully de-identified demonstration data (MIMIC-IV Demo). No actual patient data was used.

IRB Status: Not applicable. Institutional Review Board approval was not required as no human subjects research was conducted per 45 C.F.R. § 46.102(l)(2)(i) - research involving only de-identified publicly available information.

Data Availability: All data sources are publicly accessible: - Synthea OMOP: AWS S3 (no authentication required) - MIMIC-IV Demo: PhysioNet (open access) - Code and analysis scripts: <https://github.com/watilde/Harmonia>

DATA AVAILABILITY

Complete framework code and experimental data: <https://github.com/watilde/Harmonia/tree/main/research/identification-sensitivity-adaptation>

Individual module repositories:

- Module 2: [.../1-federated-partial-identification](#)
- Module 3: [.../2-federated-evalues](#)
- Module 4: [.../3-design-failure-aware-causal](#)

Synthea generator: <https://synthetichealth.github.io/synthea/>

End of Manuscript v1.0 (Revised)